

Different Model providers and their models released in last 12 months

1. Google DeepMind

- **Model:** Gemini 3.1 Pro
- **Release Date:** February 12, 2026
- **Knowledge Cutoff:** January 2025
- **Cost of Tokens:** \$1.25/1M input | \$5.00/1M output
- **Context Window:** 2,000,000 tokens
- **Availability:** Public model (Vertex AI & AI Studio)
- **When to use:**
 - Works well for large PDFs, datasets, and logs
 - Strong in video + audio + text combined reasoning
 - Useful for research-heavy testing scenarios and data validation

2. Anthropic

- **Model:** Claude Opus 4.7
- **Release Date:** April 16, 2026
- **Knowledge Cutoff:** May 2025
- **Cost of Tokens:** \$5.00/1M input | \$25.00/1M output
- **Context Window:** 1,000,000 tokens
- **Availability:** Public model
- **When to use:**
 - Accurate image understanding
 - Deep thinking – Complex architecture decisions – when system is not simple like Multiple services, APIs, Databases
 - Double-checking answers – reduces mistakes in complex tasks

3. Meta

- **Model:** Muse Spark

- **Release Date:** April 10, 2026
- **Knowledge Cutoff:** Late 2025
- **Cost of Tokens:** Integrated into Meta AI platforms
- **Context Window:** N/A (Focus on specialized modes)
- **Availability:** Closed proprietary model (First from Meta Superintelligence Lab)
- **When to use:** Proactive multimodal agents for visual comprehension, shopping, and personalized assistance.
 - Proactive multimodal agents
 - ❖ Multimodal → understands text + images + video (sometimes audio) together
 - ❖ Agent → can take actions, not just answer
 - ❖ Proactive → doesn't wait for you → suggests things on its own
 - Visual comprehension
 - Shopping assistance
 - Personalized assistance

4. Mistral AI

- **Model:** Mistral Medium 3.5
- **Release Date:** April 2026
- **Knowledge Cutoff:** Early 2026
- **Cost of Tokens:** Competitive enterprise tier
- **Context Window:** 128,000 tokens
- **Availability:** Public model
- **When to use:**
 - Efficient enterprise-grade reasoning
 - ❖ Efficient → gives good performance at lower cost & faster speed
 - ❖ Enterprise-grade → reliable enough for real business applications
 - ❖ Reasoning → can handle logic, decisions, structured thinking
 - This model fits well in step-by-step processes
 - Can be self-hosted / deployed in private cloud
 - Better control over sensitive data
 - No need to send everything to external APIs

5. DeepSeek

- **Model:** DeepSeek-V4 Pro

- **Release Date:** April 25, 2026
- **Knowledge Cutoff:** Late 2025
- **Cost of Tokens:** \$0.14/1M input | \$0.28/1M output
- **Context Window:** 1,000,000 tokens
- **Availability:** Public model (API & Web)
- **When to use:** Extreme cost-efficiency for high-performance agentic engineering and coding
 - Delivers strong performance at a much lower cost than many top-tier models
 - Strong in coding, debugging, and logic
 - Handles technical reasoning tasks well
 - Good for developer-heavy workflows
 - Building AI agents that can act, not just respond

An agent can:

- ❖ Plan steps
- ❖ Call APIs/tools
- ❖ Execute tasks
- ❖ Re-evaluate results

6. xAI

- **Model:** Grok 4.3
- **Release Date:** May 1, 2026
- **Knowledge Cutoff:** December 2025
- **Cost of Tokens:** Exclusive to X Premium+ subscribers
- **Context Window:** 128,000 tokens
- **Availability:** Public model (X platform)
- **When to use:**
 - But this model can access live information especially from social platforms. It doesn't just *remember*, it can look up what's happening now
 - Read thousands of posts quickly, understand what people are saying and Identify patterns, sentiment, and trends

7. Zhipu AI

- **Model:** GLM-5.1

- **Release Date:** May 2026
- **Cut-off Date:** Late 2025
- **Cost of Tokens:** ~1/3 of Western frontier costs
- **Context Window:** 200,000 tokens
- **Availability:** Open-weight (MIT)
- **When to use:**
 - often available as open-weight / flexible deployment
 - very good at writing, understanding, and debugging code
 - it handles calculations + reasoning + proofs
 - Can be customized or fine-tuned
 - Easier to use in internal/private systems
 - Lower cost compared to closed models

8. Cohere

- **Model:** Command R3
- **Release Date:** Late 2025
- **Knowledge Cutoff:** Mid-2025
- **Cost of Tokens:** Varies by cloud provider
- **Context Window:** 128,000 tokens
- **Availability:** Public model (Enterprise focus)
- **When to use:** High-precision RAG (Retrieval-Augmented Generation) and enterprise automation.
 - Retrieval → fetches data from your own sources (DB, docs, APIs)
 - Generation → creates answers using that data
 - High-precision → gives accurate, grounded answers (less hallucination)
 - Integrate with databases, APIs, tools
 - Automate repetitive processes
 - Handle structured workflows

9. Microsoft

- **Model:** Phi-4.1 (MoE)
- **Release Date:** March 2026

- **Knowledge Cutoff:** Late 2025
- **Cost of Tokens:** Free (Open-weight) / API via Azure
- **Context Window:** 128,000 tokens
- **Availability:** Open-weight model
- **When to use:**
 - Designed to run well even when you don't have powerful GPUs or cloud infra
 - Delivers strong reasoning and coding for its size
 - "Edge" = close to the user, not in the cloud
 - It uses multiple smaller "experts", Each expert handles a specific task and only relevant experts activate → saves compute
 - Low-resource local environments

10. Snowflake

- **Model:** Arctic 2.0
- **Release Date:** April 2026
- **Knowledge Cutoff:** Late 2025
- **Cost of Tokens:** Optimized for Snowflake Cortex
- **Context Window:** 128,000 tokens
- **Availability:** Open-weight model
- **When to use:**
 - Works well with data warehouses, analytics platforms, BI tools
 - Understands structured data (tables, schemas, pipelines)
 - Helps turn raw data → insights
 - Convert natural language → SQL queries
 - Generate optimized queries
 - Understand complex joins, filters, aggregations